

**BIRLA INSTITUTE OF TECHNOLOGY & SCIENCE, PILANI (RAJ)**  
**I SEMESTER 2025-2026**  
**ASSIGNMENT**

**Course No.: CS/SS G527**  
**Deadline: As on Nalanda**

**Course Title: Cloud Computing**  
**Maximum Marks: 60 (15%)**

---

**Note:**

- **Group of maximum four students**
  - **Submit deliverables on Nalanda (only one per group)**
  - **Demonstration may be conducted after submissions.**
- 

In this problem, you will implement a cloud application with at least 6 microservices using (multiple) cloud provider environment, out of which at least one being a web service accessible over public URL.

Following requirements must be met.

- a. All cloud infrastructure (networks, clusters, databases, storage, IAM policies, etc.) across all cloud providers must be provisioned exclusively using an Infrastructure as Code (IaC) platform (e.g., Terraform).
- b. Choose any domain such as education, healthcare etc. Your application must consist of at least six microservices in that domain, each serving a distinct functional purpose. Data analytical service (1 out of 6) must run on a different cloud provider (B). Include at least one serverless function (e.g., AWS Lambda, Google Cloud Functions, or Azure Functions) to perform an asynchronous, event-driven task. Microservices must communicate using cloud-native mechanisms such as REST, gRPC, or message queues (e.g., Kafka, AWS SQS, Google Pub/Sub).
- c. Use a managed K8s service (e.g., EKS / GKE / AKS) to host stateless microservices. Implement Horizontal Pod Autoscalers (HPAs) for at least two critical services, configured to scale based on CPU or memory utilization.
- d. All Kubernetes deployments and application updates must be managed via a GitOps controller (e.g., ArgoCD, Flux). The controller must track a Git repository. Direct kubectl apply is forbidden for service deployment.
- e. One of your services must be a real-time stream processing service (Flink) running on the managed cluster (e.g., Google Dataproc, AWS EMR, or Azure Synapse/HDInsight) in Provider B. This service must consume event from a Kafka topic. It must perform a stateful, time-windowed aggregation (e.g., "count of unique users per 1-minute window"). The

aggregated results must be published back to a separate 'results' Kafka topic. The Kafka cluster itself must be a managed cloud service (e.g., MSK, Confluent Cloud, Aiven).

- f. Use at least one distinct cloud storage products. For example: an object store (e.g., S3, GCS) for raw data (e.g., file uploads that trigger the serverless function), a managed SQL database (e.g., RDS, Cloud SQL) for relational data (e.g., user accounts, structured metadata), a managed NoSQL database (e.g., DynamoDB, Firestore, MongoDB Atlas) for high-throughput, semi-structured data (e.g., session state, real-time analytics results).
- g. Implement a comprehensive observability stack. Metrics: Deploy Prometheus and Grafana. Create a dashboard showing key service metrics (RPS, error rate, latency) and Kubernetes cluster health. Logging: Implement a centralized logging solution (e.g., EFK stack, Loki, or Cloud Logging) to aggregate logs from all microservices, including the analytics job.
- h. Use a load testing tool (e.g., k6, JMeter, Gatling) to validate your system's resilience and scalability. Load Test: Generate sustained traffic to demonstrate the HPA scaling out your services.

Deliverables:

- Design document containing System overview, Cloud deployment architecture, Microservices architecture digrams, microservice responsibilities, interconnection mechanisms, and the rationale behind design choices in the cloud.
- Microservices code (in Github repository), all IaC script code, all Kubernetes manifests and GitOps configuration
- Each student, for their part, should record a video with his/her own voice explaining the different sections in source/script code or configuration they have written with <idno> visible in the terminal or code. Video link must be present in <idno>\_video.txt. Based on video, marks will be allotted to individuals. 12M.
- A video demonstrating end to end working of your application including the testing phase must be recorded with explanation. Video link must be present in demo\_video.txt

[60M]

--&--